



TECH EXPO 2026

Track A: AI Architectures & Edge Systems

International Symposium

Nexa-Core: Next-Generation Neuromorphic Architecture

Sub-Milliwatt Spiking Neural Network Processing with Dynamic Event-Driven Synapse Fabric

Dr. Marcus Vance¹Dr. Elena Rostova¹Sarah Chen²David K. Miller³¹Nexa AI Research Institute, San Francisco, CA²Synapse Systems Laboratory, Zurich, Switzerland³Vertex Quantum & Neuromorphic Center, Austin, TX

Contact & Demo:

vance@nexa-ai.org

Portal:

https://nexa-ai.org/core

Booth: #A-104



Introduction & Motivation

Standard von Neumann architectures and traditional edge GPUs suffer from memory wall bottlenecks and high idle power dissipation during real-time event-stream processing. Edge computing in autonomous robotics and biomedical implants demands ultra-low power envelopes below 1 mW.

The Nexa-Core Paradigm:

- ✓ **Event-Driven Computation:** Operations execute sparsely only upon receiving dynamic Address Event Representation (AER) spikes.
- ✓ **In-Memory Synaptic Computing:** Non-volatile memristive crossbars perform analog matrix-vector multiplications directly inside memory arrays.
- ✓ **Sub-Millisecond Inference:** Asynchronous Leaky Integrate-and-Fire (LIF) neurons deliver ultra-fast temporal classification.

Key Architectural Breakthroughs:

0.85 mW Active Power | 0.42 ms Latency | 12.4× Energy Efficiency Boost



Microarchitecture & Synapse Fabric

The Nexa-Core IC consists of a distributed grid of Neuromorphic Tile Arrays connected via an Asynchronous Network-on-Chip (NoC). Each tile integrates an AER Event Router, a 512 × 512 Memristive Synaptic Crossbar, and 256 LIF Digital Neuron Cores.



Synaptic Crossbar Array Details: The memristive crossbar utilizes Hafnium Oxide (HfO₂) resistive switching elements, enabling 4-bit synaptic weight quantization with low-drift state retention.



Mathematical & Algorithmic Formulation

Spiking Neural Networks model membrane dynamics continuously. The internal state $V_i(t)$ of the i -th LIF neuron evolves according to the differential equation:

$$\tau_m \frac{dV_i(t)}{dt} = -(V_i(t) - V_{rest}) + R_i \sum_j W_{ij} S_j(t - \Delta_{ij}) + I_{bias}(t)$$

where τ_m represents the membrane time constant, W_{ij} is the synaptic efficacy from neuron j to i , and $S_j(t)$ denotes incoming impulse event trains.

Spike Generation Threshold Condition:

$$S_i(t) = \delta(t - t_i^*) \text{ when } V_i(t) \geq V_{th}, \quad V_i(t^*) \leftarrow V_{reset}$$

Surrogate Gradient On-Chip Backpropagation: To train SNNs end-to-end despite non-differentiable step functions, we introduce a smooth arc-tangent surrogate gradient operator:

$$\frac{\partial S_i}{\partial V_i} = \frac{1}{\pi \gamma \left[1 + \left(\frac{V_i - V_{th}}{\gamma} \right)^2 \right]}$$

This formulation minimizes quantization loss and allows on-chip online learning with zero off-line retraining overhead.



Hardware Microarchitecture Benchmarks

Fabricated in a 16nm FinFET process, Nexa-Core v2 demonstrates superior performance-per-watt metrics over existing embedded AI accelerators.

Architecture	Tech Node	Power (mW)	Latency (ms)	Energy (pJ/SOP)
Nexa-Core v2	16 nm	0.85	0.42	1.85
Vertex-SNN v1	28 nm	4.30	1.15	7.40
Apex-Edge GPU	8 nm	1250.00	8.90	142.00
Meridian TPU-v4e	7 nm	350.00	3.10	48.50

* SOP: Synaptic Operation (equivalent to MAC in ANN). Evaluated on DVS-Gesture event datasets.



Experimental Results & Performance

We benchmarked Nexa-Core against industry-standard embedded hardware platforms across key temporal classification tasks.



99.1%

DVS-Gesture Accuracy

0.42 ms

Ultra-Low Latency

1.85 pJ

Energy per SOP



Commercial Applications & Conclusion

Target Industry Deployments:

- ✓ **Nimbus-Aero Autonomous Drones:** Real-time vision-based collision avoidance under high velocity flight.
- ✓ **Synapse Medical Implants:** Ultra-low power cardiac wave anomaly detection and neural prosthetics control.
- ✓ **Meridian Smart Grid:** Microsecond high-frequency voltage fluctuation monitoring.

Summary: Nexa-Core bridges the gap between biological efficiency and artificial edge intelligence, opening new frontiers for zero-latency, sub-milliwatt autonomous computing.



References & Acknowledgments

- [1] Vance, M. et al. (2025). *Event-Driven Neuromorphic Architecture for Edge AI*. IEEE J. Solid-State Circuits, 60(4), 1102–1115.
- [2] Rostova, E. & Chen, S. (2025). *Surrogate Gradient On-Chip Backpropagation in Memristive Arrays*. Nature Electronics, 8(2), 145–156.
- [3] Miller, D. K. (2024). *Asynchronous AER Network-on-Chip Fabrics*. Proc. ISCA 2024, pp. 89–101.

Supported by Nexa AI Research Grant (2023-2026) and Vertex Systems Foundation Fellowship.